

ECON484 Machine Learning

1. Concepts of Machine Learning (2/2)

Lecturer:

Bora GÜNGÖREN


bora.gungoren@atilim.edu.tr

An Applied Example

- We will look at the basic problem of polynomial curve fitting (adapted from Bishop 2006).
 - This will involve optimizing the parameters (polynomial coefficients) and be classified as an estimation problem.
 -

An Applied Example

- Our actual function is a sine wave $t = \sin(2\pi x)$.
- We are given a training set $X = (x_1, x_2, \dots, x_N)^T$
- The data set has generated data and some (Gaussian) noise.




An Applied Example

- We should have a model and the model parameters.
- $\mathbf{w}=(w_0, w_1, w_2, \dots, w_M)$ are the model parameters where M is the dimension of the polynomial.

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j \quad (1.1)$$


An Applied Example

- Our parameter estimation problem is now expressed as “estimate \mathbf{w} (as vector) to minimize”.
- Error is defined as $E(\mathbf{w})$ which is calculated using estimated y_N and actual t_N in the training set.



$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 \quad (1.2)$$

An Applied Example



An Applied Example

- $M=3$ has more error in the training data set than $M=9$, however it will give a better overall fit for the curve.
 - This is why actual performance should be measured with practical application.
 - The problem with $M=9$ is what we call **over-fitting**.


An Applied Example

- To measure performance, we generate a new test data set (again by adding noise) of size 100.
 - Then we define the error measure as E_{RMS} .
 - Let's calculate E_{RMS}

$$E_{\text{RMS}} = \sqrt{2E(\mathbf{w}^*)/N} \quad (1.3)$$

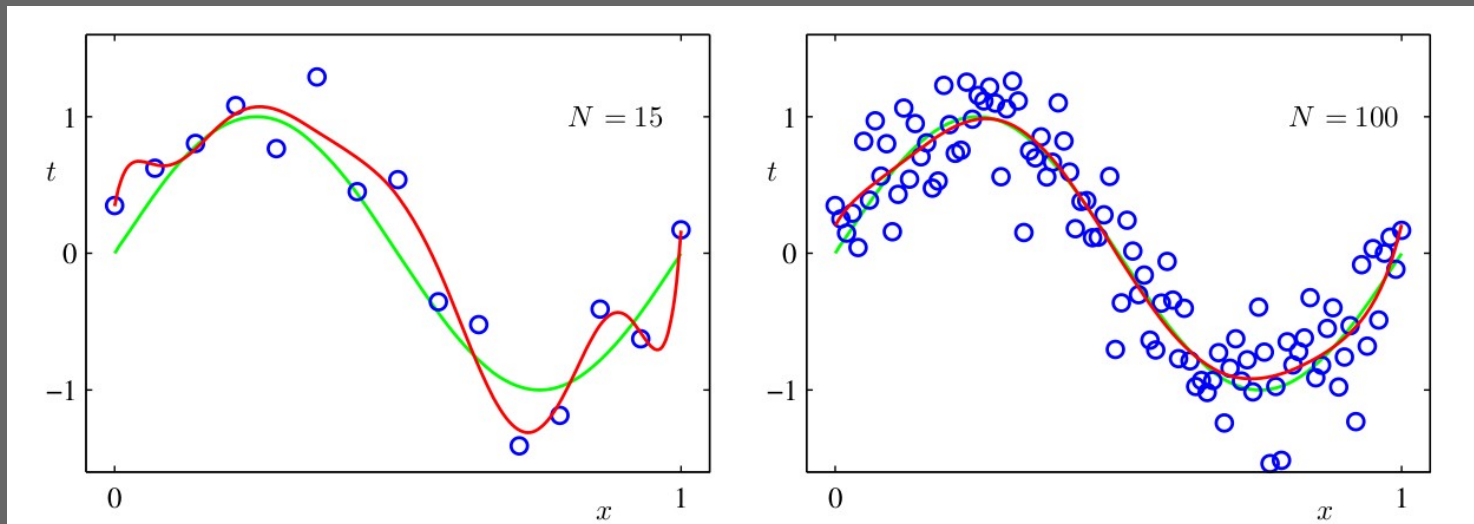
An Applied Example

- What we see
 - In the overtrained model, E_{RMS} increases wildly with test data set.
 - From $M=3$ to $M=8$ the error looks minimal.
 - So we keep the simplest form, $M=3$.



An Applied Example

- A way to reduce the error with $M=9$ model would be to train with a larger training data set.



An Applied Example

- It seems that having a larger training data set solves some problems.
 - But it is not practical to assume we have a larger training data set in all practical settings.
 - One approach is to retrain our data models frequently.
 - Another approach is to employ different models.
- Model complexity should reflect the complexity of the real world.

An Applied Example

- In the example case, the reason for the large fluctuation of estimated values is the very large coefficients in the model we obtain after training.
 - The better models include smaller coefficients.
 - So why not penalize very large coefficients in the model?


An Applied Example

- We penalize very large coefficients through a modified error definition.
 - This modification type is called a **shrinkage** (or **weight decay**), and this particular one is called a **ridge regression**.

$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2 \quad (1.4)$$

An Applied Example

- Our selection of model modification has added a new parameter to optimize
 - **Lambda** which symbolized how important the shrinkage is.
 - We should also select this parameter.



An Applied Example

- Another approach to the same estimation problem would be Bayesian.
 - Let our training data $D=(t_1, t_2, \dots, t_N)$ represent a number of observations.
 - Let us assume our parameter set \mathbf{w} is now probabilistic with a probability distribution $\mathbf{p}(\mathbf{w})$.
 - We try define all quantities as a function of \mathbf{w} .

An Applied Example

- With $p(\mathcal{D}|\mathbf{w})$ as the likelihood function:

$$p(\mathbf{w}|\mathcal{D}) = \frac{p(\mathcal{D}|\mathbf{w})p(\mathbf{w})}{p(\mathcal{D})} \quad (1.43)$$

$$p(\mathcal{D}) = \int p(\mathcal{D}|\mathbf{w})p(\mathbf{w}) \, d\mathbf{w}. \quad (1.45)$$

An Applied Example

- With the frequentist school of thought, we usually use a maximum likelihood estimator, to maximize the likelihood function $p(D|w)$.
 - The error is defined as the negative logarithm of likelihood function. As one increases the other should decrease.
 - However the fundamental difference is that we see w as a constant and discuss the variation of D given a w .

An Applied Example

- With the Bayesian school of thought, there is only the observations D , and we evaluate the uncertainty in w , after we observed D .
 - Therefore D is a constant, and there is a variation in w , which we discuss (and optimize).
- Also note that the likelihood function is not a probability distribution.
- The Bayesian approach requires a lot more computation than the frequentist approach, and is often infeasible to compute by hand.
 - But with computer resources at hand, it has become more popular.

An Applied Example


- So how to approach curve fitting from a Bayesian perspective?
 - We treat the observations $t=(t_1, t_2, \dots, t_N)^T$ as having some uncertainty with a selected probability density function (pdf). Our **model selection** is now based on our choice of PDF.
 - This is a **very big** issue. Frequentists often criticize Bayesian approach saying “selection of PDF after observation of training data is cheating”.
 - Let's assume that the distribution of error is Gaussian (which we know to be true in this case).

An Applied Example

- With this model selection any observation t has now a conditional probability based on x with the parameters \mathbf{w} , and β .
 - We use the training data set and a **maximum likelihood estimator to find the parameter values**.
 - In order to minimize error, defined as negative logarithm of the maximum likelihood function, we will try to **maximize the logarithm of the maximum likelihood function**, with respect to β .
 - Once we have the parameters, we can start making predictions.

An Applied Example

$$p(t|x, \mathbf{w}, \beta) = \mathcal{N}(t|y(x, \mathbf{w}), \beta^{-1}) \quad (1.60)$$



An Applied Example

$$p(\mathbf{t}|\mathbf{x}, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(t_n | y(x_n, \mathbf{w}), \beta^{-1}). \quad (1.61)$$

$$\ln p(\mathbf{t}|\mathbf{x}, \mathbf{w}, \beta) = -\frac{\beta}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{N}{2} \ln \beta - \frac{N}{2} \ln(2\pi). \quad (1.62)$$

An Applied Example

$$\frac{1}{\beta_{\text{ML}}} = \frac{1}{N} \sum_{n=1}^N \{y(x_n, \mathbf{w}_{\text{ML}}) - t_n\}^2. \quad (1.63)$$

- Having determined w and β , we can make predictions.
 - What we have is a predictive distribution $p(t \mid x, w, \beta)$ which gives **a probability distribution over t , not a point estimate.**
 - We could also assume w is probabilistic, and have **hyperparameters** on the distribution of w . For example **$p(w \mid \alpha)$.**

An Applied Example


- When we assume w is probabilistic, we estimate the **most likely value of w given the observations D** .
 - This is equivalent to minimizing the error function to find w .
 - Note that in Bayesian approach we take integrals which is computationally expensive.
- In the maximum likelihood approach we have demonstrated the importance of overfitting, and having a large training data set.
- In the Bayesian approach we have demonstrated how complex it could become, and how much computation intensive our tasks can become.

How to work with a small training data set?

- This is an essential question.
 - When there is a large training data set, we select some for training and some for validation.
 - When there is a small training data set, we employ **cross-validation**.
 - In this approach, we have multiple runs, each time with a different part of the training set reserved for validation.
 - We develop descriptive statistics (ie. averages) from the performance of these runs and evaluate the performance based on these descriptives.

How to work with a small training data set?

- S-fold cross-validation
 - We create S (equal sized) parts of the training data set.
 - We train using the majority (white) and validate with the withheld (red) data set.
 - We take the average of performance values.
- Extreme case is to have $S=N$ runs, which is also called **leave-one-out**.





Dimensionality

- In our simple example we discussed how we develop a model to **estimate a function of a single variable**.
 - When we have multiple variables (higher dimensionality) this becomes much more complex.

Dimensionality

- We will again use the example from Bishop 2006.
 - The example involves focusing on two variables (out of 12) in a classification problem.
 - The proposed technique divides the space into “zones” which belong to the majority class in the test data set.
 - If a point falls into a zone, then it is classified as from that class.

Dimensionality



Dimensionality

- This approach works “OK” for 3 classes and a 2-dimensional space.
 - How about more classes and the full 12-dimensional space?
 - How many zones do we have? 2^{12} ? 3^{12} ? 4^{12} ?
 - How many coefficients will we need? 12^2 ? 12^3 ? 12^4 ?
- Naive approaches just don't work with higher dimensions.
 - This is called the **curse of dimensionality** (Bellman, 1961).

Decision Theory and Machine Learning

- Decision theory proposed many methods to make decisions. They usually focus on **optimizing some parameters** in order to:
 - Minimizing false decisions
 - Minimizing expected losses
 - Detecting hard to make decisions

Decision Theory and Machine Learning

- Also consider von Neumann's minimax theorem and Yao's principle with multiple turn decision making in a competitive environment.
 - The observations \mathbf{D} we have on the competitor's decisions under partially observed environment (test data set) can be used to learn about (ie. estimate a range on) the parameters (\mathbf{w}), hence **the range of decisions provided by a particular decision making model** we assume she is using.
 - Model selection, parameters as probability distributions, observations as given facts...
 - These are the cornerstones of a Bayesian approach.
 - Bayesian Decision Theory uses machine learning in practical applications.

Model Evaluation Criteria

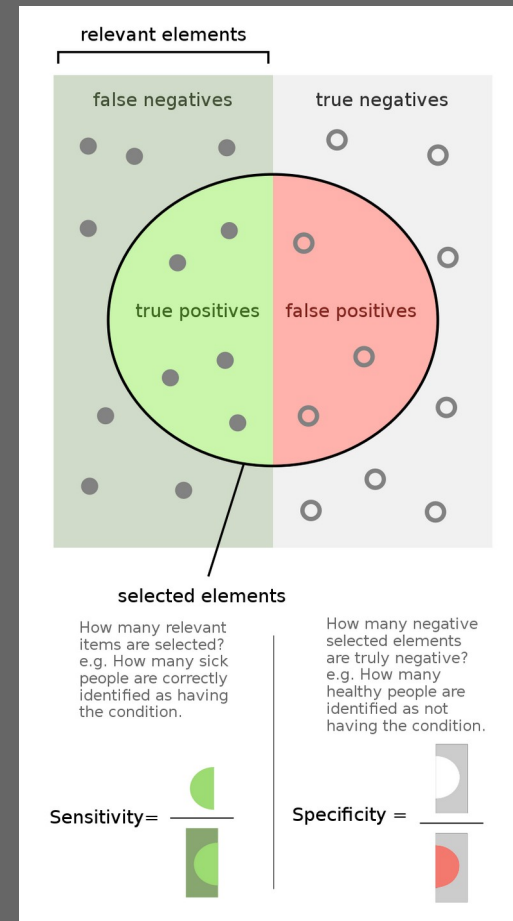
- As we see, we usually design our model prior to training it with the training data set, and then validate it with the validation data set.
 - We should be able to observe problems with the model before we use it in practice.
- Hence we need some model evaluation criteria.

Model Evaluation Criteria

- In classification problems we will have multiple type of problems.
 - The core performance measure, **classification accuracy** is the ratio of correct predictions (true positives and true negatives) to total predictions made. However in many cases, **error rate** as the rate of all other cases (false negatives, false positives) is also important.
 - In the two-class problem we define the true positive, false positive, true negative, false negative approach easily.

Model Evaluation Criteria

- Sensitivity (recall) and specificity are two important performance criteria we use. (Figure from is from Wikipedia).
- High sensitivity is always good.
- High specificity may be more important than high sensitivity because in many applications false positives cause hard to repair damages.



Model Evaluation Criteria

- A confusion matrix is a technique for summarizing the performance of a classification algorithm.
 - We have a series of classifications, actual and predicted and we see their true/false positive/negative nature.
 - We summarize using a matrix.
 - We calculate precision, accuracy and recall.

		Predicted	
		Negative (N) -	Positive (P) +
Actual	Negative -	True Negatives (TN)	False Positives (FP) Type I error
	Positive +	False Negatives (FN) Type II error	True Positives (TP)

Model Evaluation Criteria

		CONDITION determined by "Gold Standard"			
TOTAL POPULATION		CONDITION POS	CONDITION NEG	PREVALENCE $\frac{\text{CONDITION POS}}{\text{TOTAL POPULATION}}$	
TEST OUT-COME	TEST POS	True Pos TP	Type I Error False Pos FP	Precision Pos Predictive Value $\text{PPV} = \frac{\text{TP}}{\text{TEST P}}$	False Discovery Rate $\text{FDR} = \frac{\text{FP}}{\text{TEST P}}$
	TEST NEG	Type II Error False Neg FN	True Neg TN	False Omission Rate $\text{FOR} = \frac{\text{FN}}{\text{TEST N}}$	Neg Predictive Value $\text{NPV} = \frac{\text{TN}}{\text{TEST N}}$
ACCURACY ACC $\text{ACC} = \frac{\text{TP} + \text{TN}}{\text{TOT POP}}$		Sensitivity (SN), Recall Total Pos Rate TPR $\text{TPR} = \frac{\text{TP}}{\text{CONDITION POS}}$	Fall-Out False Pos Rate FPR $\text{FPR} = \frac{\text{FP}}{\text{CONDITION NEG}}$	Pos Likelihood Ratio LR + $\text{LR} + = \frac{\text{TPR}}{\text{FPR}}$	Diagnostic Odds Ratio DOR $\text{DOR} = \frac{\text{LR} +}{\text{LR} -}$
		Miss Rate False Neg Rate FNR $\text{FNR} = \frac{\text{FN}}{\text{CONDITION POS}}$	Specificity (SPC) True Neg Rate TNR $\text{TNR} = \frac{\text{TN}}{\text{CONDITION NEG}}$	Neg Likelihood Ratio LR - $\text{LR} - = \frac{\text{TNR}}{\text{FNR}}$	

Model Evaluation Criteria

- When we have more than 2 classes, calculating TP, TN, etc becomes more complex, but not more difficult.

Model Evaluation Criteria

- R^2 (coefficient of determination) is the proportion of the variation in the dependent variable that is predictable from the independent variable(s).
 - It is used to evaluate the performance of regression based estimation models.
 - An R^2 of 1 indicates that the regression predictions perfectly fit the data.
 - However, a high R^2 can occur in the presence of mis-specification of the functional form of a relationship or in the presence of outliers that distort the true relationship.
 - In many cases SSE (sum of squared errors) is preferred over R^2 .

Homework Assignment #3

- Construct a Confusion Matrix based on the classifications made by a hypothetical classifier.
 - You will use the data set that will be provided in the course Github site. The data set has classification results in the validation phase of a 5-class classifier.
 - You will do this by developing code.
 - Submissions in R, Python, Java, etc are accepted.
 - You can also submit a Notebook file.
- Due Date: March 28th
- Submission: Through Github.

Questions?

CONTACT:

bora.gungoren@atilim.edu.tr

License: Creative Commons Attribution Non-Commercial Share Alike 4.0 International (CC BY-NC-SA 4.0)